

Path-based SST relies on classical results from labelled combinatorial structures [11] to uniformly sample the set of program paths with length in $[1, \ell]$. Each path sample is provided to a constraint solver (oracle) and labelled as feasible or infeasible; see [9] and references therein. The infeasibility of a given path arises if it violates some dependencies between different parts of the program, referred to as *XOR patterns*. For instance if two if nodes are based on an unchanged expression, then their successors are correlated in every feasible path (if the program path includes the then successor of the first if node, it must also include the then successor of the second if node).

Because of the small number of available labelled paths (due to the labelling cost) compared to the complexity of the “natural” search space, i.e. that of long strings on a large alphabet, a frugal propositional representation inspired by Parikh maps [12] is considered. For $\ell \in \{1, 0, 0, 0, 0\}$, let $s[\ell]$ denote the ℓ -th symbol in string s , set to value v_ℓ if the length of s is less than ℓ .

- To each symbol v , is associated an integer attribute a_v ; $a_v(s)$ is the number of occurrences of symbol v in path s .
- To the i -th occurrence of a symbol v , is associated a categorical attribute $a_{v,i}$. Attribute $a_{v,i}(s)$ gives the next ‘informative’ symbol following the i -th occurrence of symbol v in s (or v_ℓ if s contains less than i occurrences of v).

Preliminary attempts at discriminant learning have been hindered by the tiny percentage of the feasible paths, as could have been expected from [8]. A generative learning approach was then considered.

3 Overview of EXIST

This section describes a sampling algorithm called *EXIST* for *Exploration vs eXploitation Inference for Software Testing*, able to retrieve distinct feasible paths with high probability based on a set \mathcal{E} of feasible/infeasible paths. \mathcal{E} , initially set to a small set of labelled paths, is gradually enriched with the paths generated by *EXIST* and labelled by the constraint solver.

EXIST proceeds by iteratively exploiting and updating a probabilistic model \mathcal{P} . *EXIST* involves two modules: the *Init* module estimates the probability for a path to be feasible conditionally to its extended Parikh description²; the *Decision* module uses the \mathcal{P} model to iteratively construct the current path s .

3.1 Decision module

Let s (resp. v) denote the path under construction (resp. the last node symbol in s). Let i be the total number of occurrences of v in s . Let w be one possible successor node of v ; if w is selected, the total number of w symbols in the final path will be at least the current number of occurrences of w in s , plus one; let η_w denote this number.

Let us define $p_s(w)$ as the probability for a path S to be feasible conditionally to $E_{s,w}(S) \equiv [a_{v,i}(S) \equiv w] \wedge [a_w(S) \geq \eta_w]$, estimated by the *Init* module; $p_s(w)$ is conventionally set to 0 if there is no path in \mathcal{E} satisfying $E_{s,w}$.

Probabilities $p_s(w)$ for w ranging over the successors of v are used to select the next node in s . Three options have been considered in order to favor the generation of a new feasible path.

The *Greedy* option selects the successor node w maximising $p_s(w)$.

The *RouletteWheel* option stochastically selects node w with probability proportional to $p_s(w)$.

The *BanditST* option considers the multi-armed bandit problem where every bandit arm corresponds to a successor w of the current node v and the associated reward is $p_s(w)$, and uses the UCB1 algorithm [1] for determining the best arm/successor node.

3.2 Init module

The *Init* module determines how the conditional probabilities used by the *Decision* module are estimated. The baseline *Init* option computes $p_s(w)$ as the fraction of paths in \mathcal{E} satisfying $E_{s,w}$ that are feasible. However, this option fails to guide *EXIST* efficiently due to the disjunctive nature of the target concept, as shown on the following toy problem.

²Formally, $a_{v,i}(s)$ is set to $s[t(i)] \in \mathcal{V}$, where $t(i)$ is the index of the i -th occurrence of symbol v in s ; i is initially set to 1; in case $a_{v,i}$ takes on a constant value over all examples, i is incremented.

³This probabilistic model space is meant to avoid the limitations of probabilistic ESAs and Variable Order Markov Models [24]. On one hand, probabilistic ESAs (and likewise simple Markov models) cannot model the long range dependencies of the *XOR patterns*. On the other hand, although Variable Order Markov Models can accommodate such dependencies, they are ill-suited to the sparsity of the initial data available.